

# PGIMER\_Assignment\_Rohit\_Final

May 30, 2025

```
[1]: # =====  
# Author: Rohit Gupta  
# Contact : Email : rohit.gupta.delhi1995@gmail.com  
# PGIMER Project Scientist 1 : Assignment:  
# =====  
# Import all neccesary libraries for model, training, testing, etc.  
  
#directory manipulation:  
import os  
  
# deep learning associated libraries:  
import torch  
import torchvision  
import torch.nn as nn  
import torch.optim as optim  
from torchvision import datasets, transforms, models  
from torch.utils.data import DataLoader, random_split  
  
# determine results of training,testing:  
from sklearn.metrics import classification_report, confusion_matrix  
import numpy as np  
  
# matplotlib, seaborn visualization libraries:  
import matplotlib.pyplot as plt
```

```
[2]: import os  
os.getcwd()  
#os.chdir('/Users/rohitgupta/Downloads/pgimer_assignment/')  

```

```
[2]: '/Users/rohitgupta/Documents/ml_projects'
```

```
[ ]: # Post the analysis, and retrospection, here are the answers for the questions_  
      ↪asked in the assignment.pdf:  
  
# a. Three Appropriate Metrics and Justifications:  
# The evaluation metrics I used were : Precision, Recall, F1 score, AUC-ROC_  
      ↪(curve).  
# Each of these metrics were used for the following primary reason:
```

# Precision and Recall: Crucial for medical diagnoses, high recall ensures  
↳ pneumonia cases are not ignored,

# while precision minimizes false alarms e.g Normal cases being treated as  
↳ Pneumonia inflicted).

# F1-Score: Ensures there is a balance between precision and recall values,  
↳ since both the false positive and

# false negatives have clinical consequences and economic repercussions.

# AUC-ROC: It provides a visual depiction as well as provides a quantitative  
↳ assessment for

# the model's ability to distinguish between normal and pneumonia cases across  
↳ different threshold values.

# b. Class Imbalance Detection and Mitigation:

# Class imbalance detected post running the model on the test dataset and then  
↳ seeing the model performance

# evaluation metrics. [388 normal vs 3,494 pneumonia cases]. Furthermore 0  
↳ false negatives showing up repeatedly,

# followed by 1.00 precision for Normal cases. Mitigation strategies applied:  
↳ Utilized weighted loss function,

# weighted random sampler, enhanced data augmentation technique (used geometric  
↳ transformations).

# c. Overfitting Prevention Measures:

# I applied data augmentation techniques, and utilized transforms such as  
↳ RandomResizedCrop, RandomHorizontalFlip,

# RandomRotation to training data. Furthermore, I also applied regularization  
↳ techniques such as L2 Weight Decay,

# Early stopping (with patience = 3, triggered at the epoch run at 4, to  
↳ prevent overtraining of the ResNet50 model).

# Fine tuned final classification layer to leverage learned features. However,  
↳ more improvement with this

# respect could be done, such as more layers could be unfrozen and utilized.  
↳ However due to paucity of time,

# restricted myself to this.

# Question : A short note on the hyperparameters used:

# I implemented and optimized different hyperparameters, especially post  
↳ realizing the class imbalance issue

# that manifested during the initial training, implementation of the ResNet 50  
↳ model. First started with a

# conservative learning rate of  $1e^{-4}$ ; batch size of 32 and Adam optimizer  
↳ with L2 regularization

# (using a  $e^{-5}$  weight decay).

```

# Post identification of the class imbalance, and to remedy the overfitting of
↳ the model,
# decided to further adjust the training related hyper-parameters by setting a
↳ maximum of 20 epochs,
# but with early stopping with patience = 3, which halted training at epoch 4
↳ itself, to prevent
# overfitting (meaning, to prevent the model learning/memorizing the training
↳ data itself too well).
# Finally, addressed the class imbalance again by adjusting CrossEntropyLoss
↳ with class weights as
# [1.7,1.0] and WeightedRandomSampler. Furthermore, did a couple of more
↳ adjustments with
# the hyperparameters as I continued to refine the model to improve performance
↳ - especially with
# class imbalance plaguing it.

```

```

[3]: import matplotlib.pyplot as plt
import os

# Create a directory to store the images
output_dir = 'output_images'
if not os.path.exists(output_dir):
    os.makedirs(output_dir)

# Your code to generate plots goes here
# For example:
for i in range(5):
    plt.figure()
    plt.plot([1, 2, 3, 4], [i, i**2, i**3, i**4])
    plt.title(f'Plot {i+1}')

    # Save the figure
    filename = f'plot_{i+1}.png'
    plt.savefig(os.path.join(output_dir, filename))
    plt.close() # Close the figure to free up memory

print(f"All images have been saved to the '{output_dir}' directory.")

```

All images have been saved to the 'output\_images' directory.

```

[4]: # Evaluation Strategy:
# Post the analysis the answers for the questions (in the assignment included
↳ here):

# a. Choose 3-appropriate metrics and justify your choices.
# b. Discuss how you detect and mitigate class imbalance in the training set.
# c. Describe measures taken to prevent over-fitting (e.g., data augmentation,

```

```
# regularization).
```

```
# answer a:
```

```
# answer b:
```

```
# answer c:
```

```
[5]: # test if the libraries are working or not:
```

```
import torch
```

```
import torchvision
```

```
print("PyTorch version:", torch.__version__)
```

```
print("Torchvision version:", torchvision.__version__)
```

```
print("CUDA available:", torch.cuda.is_available())
```

PyTorch version: 2.2.2

Torchvision version: 0.17.2

CUDA available: False

```
[6]: # running on mac m1 silicon, therefore required to test:
```

```
import torch
```

```
if torch.backends.mps.is_available():
```

```
    print("MPS is available!")
```

```
else:
```

```
    print("MPS not available.")
```

MPS is available!

```
[7]: device = (  
    torch.device("mps")  
    if torch.backends.mps.is_available()  
    else torch.device("cpu")  
)
```

```
[8]: # path definition:
```

```
os.getcwd()
```

```
os.chdir('/Users/rohitgupta/Downloads/pgimer_assignment/')  

```

```
[9]: os.listdir()
```

```
[9]: ['.DS_Store',  
      'requirements.txt',  
      'enviornment.yml',  
      'notebook_images',  
      'pneumoniarnist.npz',
```

```
'Assignment.pdf']
```

```
[10]: # loading the data and checking it:
from numpy import load

dataset = load('pneumoniamnist.npz')
lst = dataset.files

#for item in lst:
#    print(item)
#    print(data[item])
```

```
[11]: import numpy as np
data = np.load("pneumoniamnist.npz")

print("Keys:", list(data.keys()))
print("Train images shape:", data['train_images'].shape)
print("Can access index 2?", data['train_images'][2].shape)
```

```
Keys: ['train_images', 'train_labels', 'val_images', 'val_labels',
'test_images', 'test_labels']
Train images shape: (3882, 28, 28)
Can access index 2? (28, 28)
```

```
[12]: # Data Preparation so that it can be fed into the model (in this case the
↳ResNet-50)
# First step is to define the transforms

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.Grayscale(num_output_channels=3), # ResNet model expects 3
↳channels
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                          [0.229, 0.224, 0.225])
])

# Load dataset from folders
#dataset = datasets.ImageFolder(root="./data/train", transform=transform)
```

```
[13]: # Split into train and validation using the 80,20 rule;
# 80% of data for training, rest of the 20% for testing:

from torch.utils.data import Dataset
from PIL import Image

class PneumoniaDataset(Dataset):
    def __init__(self, images, labels, transform=None):
```

```

        self.images = images
        self.labels = labels.squeeze() #remove any extra dimensions that may
        ↪ have been there
        self.transform = transform

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        # Get image and label
        image = self.images[idx]
        label = self.labels[idx]

        # Convert to PIL Image for transforms
        image = Image.fromarray(image, mode='L') # Grayscale

        # Apply transforms
        if self.transform:
            image = self.transform(image)

        # we must ensure the image tensor is float32 for MPS compatibility
        image = image.float()

        return image, torch.tensor(label, dtype=torch.long)

# load the data
data = np.load("pneumoniarnist.npz")

# Create datasets using the existing splits from the file
train_dataset = PneumoniaDataset(
    data['train_images'],
    data['train_labels'],
    transform=transform
)

val_dataset = PneumoniaDataset(
    data['val_images'],
    data['val_labels'],
    transform=transform
)

test_dataset = PneumoniaDataset(
    data['test_images'],
    data['test_labels'],
    transform=transform
)

```

```

)

# Create DataLoaders
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True,
    ↪num_workers=0)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False,
    ↪num_workers=0)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False,
    ↪num_workers=0)

# Test that it works
print(f"Train dataset: {len(train_dataset)} samples")
print(f"Val dataset: {len(val_dataset)} samples")
print(f"Test dataset: {len(test_dataset)} samples")

# Test loading one batch
try:
    batch = next(iter(train_loader))
    print(f"Batch test successful: {batch[0].shape}, {batch[1].shape}")
except Exception as e:
    print(f"Error: {e}")

```

Train dataset: 3882 samples

Val dataset: 524 samples

Test dataset: 624 samples

Batch test successful: torch.Size([32, 3, 224, 224]), torch.Size([32])

```

[14]: print(data['train_images'].shape)
      print(data['test_images'].shape)
      print(data['val_images'].shape)

```

(3882, 28, 28)

(624, 28, 28)

(524, 28, 28)

```

[15]: # check how the images look:

```

```

from PIL import Image
import matplotlib.pyplot as plt
import numpy as np

fig, axes = plt.subplots(1, 3, figsize=(15, 5))

for i in range(3):
    img = data['train_images'][i] # shape: (28, 28)
    label = data['train_labels'][i][0]

```

```


# Normalize to 0-255, clip to prevent overflow
img_scaled = (img - img.min()) / (img.max() - img.min()) # scale 0-1
img_uint8 = (img_scaled * 255).astype(np.uint8)

# Convert to PIL, resize with smooth interpolation
img_pil = Image.fromarray(img_uint8)
img_resized = img_pil.resize((224, 224), resample=Image.BICUBIC)
img_np = np.array(img_resized)

# Display
axes[i].imshow(img_np, cmap='gray')
axes[i].set_title(f"Sample {i+1}\nLabel: {label} - {'Pneumonia' if label == 1
↵1 else 'Normal'}")
axes[i].axis('off')

plt.tight_layout()
plt.show()

```



```

[16]: # setting up the model, using pretrained resnet50 for the assigned task:
# model name is 'model_res'
model_res = models.resnet50(pretrained=True)

/opt/miniconda3/envs/jupyter/lib/python3.8/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/opt/miniconda3/envs/jupyter/lib/python3.8/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet50_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights.

```



```
warnings.warn(msg)
```

```
[17]: # Replace the final layer
num_fts = model_res.fc.in_features

model_res.fc = nn.Linear(num_fts, 2).to(device)

model_res = model_res.to(device)

# for verification
print(f"Model device: {next(model_res.parameters()).device}")
print(f"Conv1 weight device: {model_res.conv1.weight.device}")
```

Model device: mps:0

Conv1 weight device: mps:0

```
[18]: # print model details : helps understand the assigned weights, hyperparameters
      ↪ [useful later for hyperparameter tuning, optimization]
      # print(model_res)
```

```
[19]: # Establishing Loss, Optimizer, Scheduler:
criterion = nn.CrossEntropyLoss()
# Ensure the optimizer is created AFTER the model is fully moved
optimizer = optim.Adam(model_res.parameters(), lr=1e-4)
```

```
[20]: # Training the model:
      # creating the function:

      # train function :
def train(model, train_loader, val_loader, criterion, optimizer, epochs=5):
    for epoch in range(epochs):
        model_res.train()
        total_loss = 0
        correct = 0

        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model_res(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            total_loss += loss.item()
            _, preds = torch.max(outputs, 1)
            correct += torch.sum(preds == labels.data)
```

```

        # Change .double() to .float() for MPS compatibility
        acc = correct.float() / len(train_loader.dataset)
        print(f"Epoch {epoch+1}: Train Loss={total_loss:.4f}, Accuracy={acc:.
↪4f}")
        validate(model, val_loader)

# validation function:
def validate(model, val_loader):
    model_res.eval()
    correct = 0
    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model_res(inputs)
            _, preds = torch.max(outputs, 1)
            correct += torch.sum(preds == labels.data)
    # Change .double() to .float() for MPS compatibility
    acc = correct.float() / len(val_loader.dataset)
    #acc = correct.double() / len(val_loader.dataset)
    print(f"Validation Accuracy={acc:.4f}")

```

```

[21]: # run the training model:
      # we are calling the function above
      # epochs = 10 for the time being, can be changed as well

      train(model_res, train_loader, val_loader, criterion, optimizer, epochs=10)

```

```

Epoch 1: Train Loss=15.2182, Accuracy=0.9565
Validation Accuracy=0.9561
Epoch 2: Train Loss=5.9785, Accuracy=0.9815
Validation Accuracy=0.9618
Epoch 3: Train Loss=4.6641, Accuracy=0.9861
Validation Accuracy=0.9599
Epoch 4: Train Loss=1.9434, Accuracy=0.9946
Validation Accuracy=0.9542
Epoch 5: Train Loss=2.0726, Accuracy=0.9954
Validation Accuracy=0.9618
Epoch 6: Train Loss=0.9702, Accuracy=0.9972
Validation Accuracy=0.9427
Epoch 7: Train Loss=0.6306, Accuracy=0.9982
Validation Accuracy=0.9676
Epoch 8: Train Loss=1.8327, Accuracy=0.9951
Validation Accuracy=0.9695
Epoch 9: Train Loss=2.4910, Accuracy=0.9928
Validation Accuracy=0.9427
Epoch 10: Train Loss=1.0815, Accuracy=0.9969
Validation Accuracy=0.9504

```

```
[22]: validate(model_res, val_loader)
```

Validation Accuracy=0.9504

```
[24]: # Now finally check how the model fares with respect to test dataset:
from sklearn.metrics import classification_report

y_true = []
y_pred = []

model_res.eval()
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model_res(inputs)
        _, preds = torch.max(outputs, 1)

        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

# display the classification report lastly:
print(classification_report(y_true, y_pred, target_names=["NORMAL",
↪ "PNEUMONIA"]))
```

	precision	recall	f1-score	support
NORMAL	0.95	0.78	0.86	234
PNEUMONIA	0.88	0.98	0.93	390
accuracy			0.90	624
macro avg	0.92	0.88	0.89	624
weighted avg	0.91	0.90	0.90	624

```
[25]: # visualizing the results, by using AUC-ROC Curve [Area Under the Curve -
↪ Reciever Operating Characteristics]
```

```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Get softmax scores for class 1 (Pneumonia)
model_res.eval()
probs = []
true_labels = []

with torch.no_grad():
    for inputs, labels in test_loader:
```


```

inputs, labels = inputs.to(device), labels.to(device)
outputs = model_res(inputs)
probs.extend(torch.softmax(outputs, dim=1)[: , 1].cpu().numpy())
true_labels.extend(labels.cpu().numpy())

fpr, tpr, _ = roc_curve(true_labels, probs)
auc = roc_auc_score(true_labels, probs)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"AUC = {auc:.2f}")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```



```
[26]: # Plotting a confusion matrix, to understand true positives, false positives,
      ↪ true negatives, false negatives :
      # [correct diagnosis, pneumonia does not exist - but diagnosed, correct
      ↪ diagnosis - with absence of it, pneumonia does not exist - and not diagnosed]
      # Interpretation of the confusion matrix :

      # Plotting a confusion matrix to analyze:
      # True Positives (TP): Pneumonia exists - and correctly diagnosed
      # False Positives (FP): Pneumonia exists - wrongly diagnosed
      # True Negatives (TN): Normal case - correctly diagnosed
      # False Negatives (FN): Pneumonia exists - but not diagnosed (strong error)


      from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
      import matplotlib.pyplot as plt
      import numpy as np

      # use y_true, y_pred

      cm = confusion_matrix(y_true, y_pred)

      # display labels on the heatmap
      class_names = ["NORMAL", "PNEUMONIA"]

      # Display the matrix
      disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
      disp.plot(cmap=plt.cm.Blues)
      plt.title("Confusion Matrix")
      plt.grid(False)
      plt.show()
```



```
[27]: # Interesting observation : False Negative = 0, suggestive of an underlying
      ↪ issue
      # Let us check for 'class imbalance'
      # class imbalance (issues)
      # how to avoid overfitting errors

      labels = data['train_labels']
      unique, counts = np.unique(labels, return_counts=True)
      print("Class distribution:", dict(zip(unique, counts)))

      # Result : Normal(0) : 388, Pneumonia(1)
```

Class distribution: {0: 388, 1: 3494}

```
[28]: # Initiate remedial steps:

      class_weights = torch.tensor([1.7, 1.0]).to(device) # Tune based on step above
      criterion = nn.CrossEntropyLoss(weight=class_weights)
```

```
[29]: # Utilize weighted sampler instead of using train_loader:

      from torch.utils.data import WeightedRandomSampler
```

```

targets = train_dataset.labels
class_sample_count = np.array([len(np.where(targets == t)[0]) for t in np.
    ↪unique(targets)])
weights = 1. / class_sample_count
samples_weight = np.array([weights[t] for t in targets])

samples_weight = torch.from_numpy(samples_weight)
sampler = WeightedRandomSampler(samples_weight.type('torch.DoubleTensor'),
    ↪len(samples_weight))

train_loader = DataLoader(train_dataset, batch_size=32, sampler=sampler,
    ↪num_workers=0)

```

[30]: *# Redefine augmented transform :*

```

train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.Grayscale(num_output_channels=3),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
        [0.229, 0.224, 0.225])
])

train_dataset = PneumoniaDataset(
    data['train_images'],
    data['train_labels'],
    transform=train_transform
)

# not doing this for the validation and test datasets

```

[31]: *# update optimizer with added regularization:*

```
optimizer = optim.Adam(model_res.parameters(), lr=1e-4, weight_decay=1e-5)
```

[32]: *# Re-train the model and re-evaluate to check the model performance:*

```
train(model_res, train_loader, val_loader, criterion, optimizer, epochs=10)
```

```

Epoch 1: Train Loss=3.8799, Accuracy=0.9902
Validation Accuracy=0.9637
Epoch 2: Train Loss=2.2026, Accuracy=0.9941
Validation Accuracy=0.9809
Epoch 3: Train Loss=0.3695, Accuracy=0.9992
Validation Accuracy=0.9733
Epoch 4: Train Loss=0.1312, Accuracy=0.9997
Validation Accuracy=0.9599
Epoch 5: Train Loss=0.0288, Accuracy=1.0000

```

```

Validation Accuracy=0.9733
Epoch 6: Train Loss=0.1623, Accuracy=0.9997
Validation Accuracy=0.9542
Epoch 7: Train Loss=4.3638, Accuracy=0.9892
Validation Accuracy=0.9752
Epoch 8: Train Loss=1.1179, Accuracy=0.9967
Validation Accuracy=0.9714
Epoch 9: Train Loss=1.0280, Accuracy=0.9977
Validation Accuracy=0.9733
Epoch 10: Train Loss=1.0894, Accuracy=0.9972
Validation Accuracy=0.9504

```

```

[33]: # model validation:
      validate(model_res, val_loader)

```

```

Validation Accuracy=0.9504

```

```

[34]: # now on test dataset:
      from sklearn.metrics import classification_report

      y_true = []
      y_pred = []

      model_res.eval()
      with torch.no_grad():
          for inputs, labels in test_loader:
              inputs, labels = inputs.to(device), labels.to(device)
              outputs = model_res(inputs)
              _, preds = torch.max(outputs, 1)

              y_true.extend(labels.cpu().numpy())
              y_pred.extend(preds.cpu().numpy())

      # display the classification report lastly:
      print(classification_report(y_true, y_pred, target_names=["NORMAL",
↪ "PNEUMONIA"]))

```

	precision	recall	f1-score	support
NORMAL	0.96	0.61	0.74	234
PNEUMONIA	0.81	0.98	0.89	390
accuracy			0.84	624
macro avg	0.88	0.80	0.82	624
weighted avg	0.86	0.84	0.83	624




```
[35]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

# use y_true, y_pred

cm = confusion_matrix(y_true, y_pred)

# display labels on the heatmap
class_names = ["NORMAL", "PNEUMONIA"]

# Display the matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.grid(False)
plt.show()
```



```
[36]: # visualizing the results, by using AUC-ROC Curve [Area Under the Curve - 
↳ Reciever Operating Characteristics]

from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
```

```


# Get softmax scores for class 1 (Pneumonia)
model_res.eval()
probs = []
true_labels = []

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model_res(inputs)
        probs.extend(torch.softmax(outputs, dim=1)[: , 1].cpu().numpy())
        true_labels.extend(labels.cpu().numpy())

fpr, tpr, _ = roc_curve(true_labels, probs)
auc = roc_auc_score(true_labels, probs)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"AUC = {auc:.2f}")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```



```
[37]: # class imbalance still persisting above:
# False positives (1st quadrant) and False Negatives(third quadrant) high and
# low respectively.

# the training is still yielding higher on average accuracy, despite the added
# measures to curtail the overfitting phenomenon
# we will now add more stronger data augmentation and early stopping

from torchvision import transforms

train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224, scale=(0.8, 1.0)), # Zoom in/out
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=15),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1)),
    transforms.Grayscale(num_output_channels=3),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                        [0.229, 0.224, 0.225])
])
```

```
])
```

```
[38]: train_dataset = PneumoniaDataset(  
    data['train_images'],  
    data['train_labels'],  
    transform=train_transform  
)  
  
# keeping validation and test data set same
```

```
[39]: # define early stopping :  
class EarlyStopping:  
    def __init__(self, patience=3, min_delta=0.001):  
        self.patience = patience  
        self.min_delta = min_delta  
        self.counter = 0  
        self.best_score = None  
        self.early_stop = False  
  
    def __call__(self, val_acc):  
        if self.best_score is None:  
            self.best_score = val_acc  
        elif val_acc < self.best_score + self.min_delta:  
            self.counter += 1  
            if self.counter >= self.patience:  
                self.early_stop = True  
        else:  
            self.best_score = val_acc  
            self.counter = 0
```

```
[40]: # Modify training loop to incorporate the early stopping mechanism:  
  
def train(model, train_loader, val_loader, criterion, optimizer, epochs=10):  
    early_stopping = EarlyStopping(patience=3)  
  
    for epoch in range(epochs):  
        model.train()  
        total_loss = 0  
        correct = 0  
  
        for inputs, labels in train_loader:  
            inputs, labels = inputs.to(device), labels.to(device)  
            optimizer.zero_grad()  
            outputs = model(inputs)  
            loss = criterion(outputs, labels)  
            loss.backward()  
            optimizer.step()
```

```

        total_loss += loss.item()
        _, preds = torch.max(outputs, 1)
        correct += torch.sum(preds == labels.data)

    train_acc = correct.float() / len(train_loader.dataset)
    print(f"Epoch {epoch+1}: Train Loss={total_loss:.4f},  

    ↳Accuracy={train_acc:.4f}")

    val_acc = validate(model, val_loader)

    early_stopping(val_acc)
    if early_stopping.early_stop:
        print("Early stopping triggered.")
        break

```

[41]: *# updating the validation function to include the accuracy:*

```

def validate(model, val_loader):
    model.eval()
    correct = 0
    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            correct += torch.sum(preds == labels.data)

    acc = correct.float() / len(val_loader.dataset)
    print(f"Validation Accuracy={acc:.4f}")
    return acc

```

[42]: *# run the training:*

```
train(model_res, train_loader, val_loader, criterion, optimizer, epochs=20)
```

```

Epoch 1: Train Loss=0.8506, Accuracy=0.9974
Validation Accuracy=0.9695
Epoch 2: Train Loss=0.0747, Accuracy=1.0000
Validation Accuracy=0.9656
Epoch 3: Train Loss=0.0231, Accuracy=1.0000
Validation Accuracy=0.9695
Epoch 4: Train Loss=0.0147, Accuracy=1.0000
Validation Accuracy=0.9656
Early stopping triggered.

```

[43]: *# model validation:*

```
validate(model_res, val_loader)
```

Validation Accuracy=0.9656

```
[43]: tensor(0.9656, device='mps:0')
```

```
[44]: # model testing (run on test data finally):
# now on test dataset:
from sklearn.metrics import classification_report

y_true = []
y_pred = []

model_res.eval()
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model_res(inputs)
        _, preds = torch.max(outputs, 1)

        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

# display the classification report lastly:
print(classification_report(y_true, y_pred, target_names=["NORMAL",
↪ "PNEUMONIA"]))
```

	precision	recall	f1-score	support
NORMAL	0.99	0.52	0.68	234
PNEUMONIA	0.77	1.00	0.87	390
accuracy			0.82	624
macro avg	0.88	0.76	0.78	624
weighted avg	0.86	0.82	0.80	624

```
[45]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

# use y_true, y_pred

cm = confusion_matrix(y_true, y_pred)


# display labels on the heatmap
class_names = ["NORMAL", "PNEUMONIA"]

# Display the matrix
```

```

disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.grid(False)
plt.show()

```



```

[46]: #tSNE visualization:

# Load the required libraries:
from sklearn.manifold import TSNE
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import torch

# Step 1: Extract deep features from the test set using ResNet (remove final
↳ layer)
features = []
labels_list = []

# Remove final classification layer
feature_extractor = torch.nn.Sequential(*list(model_res.children())[:-1])

```

```

feature_extractor.eval()

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs = inputs.to(device)
        outputs = feature_extractor(inputs)  # e.g., [B, 2048, 1, 1] for
        ↪ ResNet-50
        outputs = outputs.view(outputs.size(0), -1)  # Flatten to [B, D]
        features.append(outputs.cpu())
        labels_list.extend(labels.cpu().numpy())

# Stack all features into a numpy array
features_np = torch.cat(features, dim=0).numpy()

# Standardize before t-SNE
features_np = StandardScaler().fit_transform(features_np)

# Apply t-SNE
tsne = TSNE(n_components=2, perplexity=30, random_state=42)
features_2d = tsne.fit_transform(features_np)

# Create a DataFrame and plot
df = pd.DataFrame()
df['x'] = features_2d[:, 0]
df['y'] = features_2d[:, 1]
df['label'] = ['NORMAL' if l == 0 else 'PNEUMONIA' for l in labels_list]


plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='x', y='y', hue='label', palette='Set1', alpha=0.7,
    ↪ s=50)
plt.title("t-SNE Visualization of ResNet Features (Test Set)")
plt.xlabel("t-SNE Component 1")
plt.ylabel("t-SNE Component 2")
plt.grid(True)
plt.legend(title="Class", loc="best")
plt.show()

```

/opt/miniconda3/envs/jupyter/lib/python3.8/site-packages/threadpoolctl.py:1019: RuntimeWarning: libc not found. The ctypes module in Python 3.8 is maybe too old for this OS.

```
warnings.warn(
```





```
[47]: # Import libraries
import umap
import torch
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.preprocessing import StandardScaler

# Extract features (same as for t-SNE)
features = []
labels_list = []

feature_extractor = torch.nn.Sequential(*list(model_res.children())[:-1])
feature_extractor.eval()

with torch.no_grad():
    for inputs, labels in test_loader:
        inputs = inputs.to(device)
        outputs = feature_extractor(inputs)
        outputs = outputs.view(outputs.size(0), -1) # flatten to [B, D]
```

```

        features.append(outputs.cpu())
        labels_list.extend(labels.cpu().numpy())

features_np = torch.cat(features, dim=0).numpy()

# Standardize
features_np = StandardScaler().fit_transform(features_np)

# Apply UMAP
umap_model = umap.UMAP(n_components=2, random_state=42)
umap_2d = umap_model.fit_transform(features_np)

# Plot UMAP
df_umap = pd.DataFrame()
df_umap['x'] = umap_2d[:, 0]
df_umap['y'] = umap_2d[:, 1]
df_umap['label'] = ['NORMAL' if l == 0 else 'PNEUMONIA' for l in labels_list]

plt.figure(figsize=(8, 6))
sns.scatterplot(data=df_umap, x='x', y='y', hue='label', palette='Set1',
               ↪alpha=0.7, s=50)
plt.title("UMAP Visualization of ResNet Features (Test Set)")
plt.xlabel("UMAP Dimension 1")
plt.ylabel("UMAP Dimension 2")
plt.grid(True)
plt.legend(title="Class", loc="best")
plt.show()

```

```

/opt/miniconda3/envs/jupyter/lib/python3.8/site-packages/umap/umap_.py:1952:
UserWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no seed
for parallelism.
    warn(

```

